


CHOCOLATE AND MAGIC MUSHROOMS DATASETS







Vrishab Prasanth Davey <300438343>

Best parameters
Confusion matrix
Precision,recall,F1 score

1. Chocolate(Imbalanced Dataset)

i. Original Dataset





Best parameters for Decision Tree: {'min_samples_split': 5, 'min_samples_leaf': 2, 'max_depth': 3, 'criterion': 'entropy'}

Confusion Matrix for Decision Tree:

```
[[ 0  4]
```

```
 [ 0 373]]
```

Precision: 0.99, Recall: 1.00, F1-score: 0.99

Best parameters for Random Forest: {'n_estimators': 50, 'min_samples_split': 5, 'min_samples_leaf': 2, 'max_depth': None, 'bootstrap': False}

Confusion Matrix for Random Forest:

```
[[ 0  4]
```

```
 [ 0 373]]
```

Precision: 0.99, Recall: 1.00, F1-score: 0.99

Best parameters for SVM: {'kernel': 'linear', 'gamma': 'scale', 'degree': 3, 'C': 100}

Confusion Matrix for SVM:

```
[[ 0  4]
```

```
 [ 0 373]]
```

Precision: 0.99, Recall: 1.00, F1-score: 0.99

Best parameters for Gradient Boosting: {'subsample': 0.8, 'n_estimators': 50, 'min_samples_split': 5, 'max_depth': 3, 'learning_rate': 0.01}

Confusion Matrix for Gradient Boosting:

```
[[ 0  4]
```

```
 [ 0 373]]
```

Precision: 0.99, Recall: 1.00, F1-score: 0.99

Best parameters for MLP: {'solver': 'sgd', 'learning_rate': 'constant', 'hidden_layer_sizes': (100, 100), 'alpha': 0.001, 'activation': 'relu'}

Confusion Matrix for MLP:

```
[[ 0  4]
```

```
 [ 0 373]]
```

Precision: 0.99, Recall: 1.00, F1-score: 0.99

Best parameters for k-NN: {'weights': 'uniform', 'p': 1, 'n_neighbors': 5}


Confusion Matrix for k-NN:

```
[[ 0  4]
```

```
 [ 0 373]]
```

Precision: 0.99, Recall: 1.00, F1-score: 0.99

ii. Undersampled



Decision Tree

Actual \ Predicted	0	1
0	3	1
1	218	155

Random Forest

Actual \ Predicted	0	1
0	1	3
1	176	197

SVM

Actual \ Predicted	0	1
0	2	2
1	166	207

Gradient Boosting

Actual \ Predicted	0	1
0	2	2
1	143	230

MLP

Actual \ Predicted	0	1
0	2	2
1	185	188

k-NN

Actual \ Predicted	0	1
0	1	3
1	160	213

Best parameters for Decision Tree: {'min_samples_split': 5, 'min_samples_leaf': 1, 'max_depth': 5, 'criterion': 'entropy'}

Confusion Matrix for Decision Tree:

```
[[ 3  1]
```

```
[218 155]]
```

Precision: 0.99, Recall: 0.42, F1-score: 0.59

Best parameters for Random Forest: {'n_estimators': 200, 'min_samples_split': 2, 'min_samples_leaf': 4, 'max_depth': 10, 'bootstrap': False}

Confusion Matrix for Random Forest:

```
[[ 1  3]
```

```
[176 197]]
```

Precision: 0.98, Recall: 0.53, F1-score: 0.69

Best parameters for SVM: {'kernel': 'rbf', 'gamma': 'auto', 'degree': 2, 'C': 100}

Confusion Matrix for SVM:

```
[[ 2  2]
```

```
[166 207]]
```

Precision: 0.99, Recall: 0.55, F1-score: 0.71

Best parameters for Gradient Boosting: {'subsample': 1.0, 'n_estimators': 200, 'min_samples_split': 2, 'max_depth': 5, 'learning_rate': 0.1}

Confusion Matrix for Gradient Boosting:

```
[[ 2  2]
```

```
[143 230]]
```

Precision: 0.99, Recall: 0.62, F1-score: 0.76

Best parameters for MLP: {'solver': 'adam', 'learning_rate': 'constant', 'hidden_layer_sizes': (50,), 'alpha': 0.0001, 'activation': 'tanh'}

Confusion Matrix for MLP:

```
[[ 2  2]
```

```
[185 188]]
```

Precision: 0.99, Recall: 0.50, F1-score: 0.67

Best parameters for k-NN: {'weights': 'distance', 'p': 2, 'n_neighbors': 3}

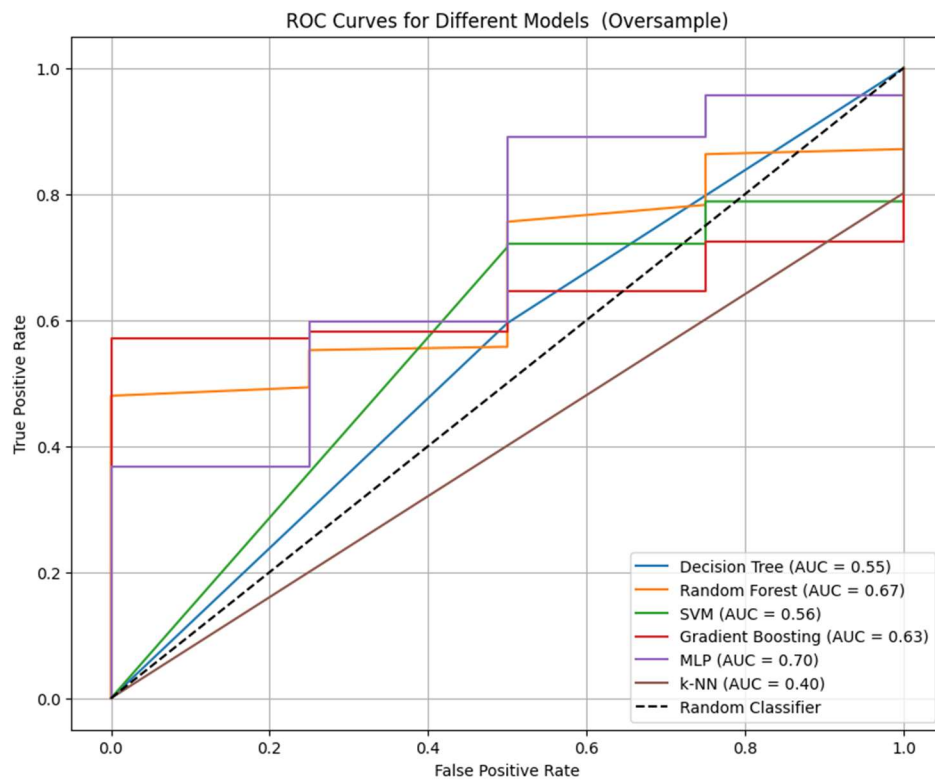
Confusion Matrix for k-NN:







$\begin{bmatrix} 1 & 3 \end{bmatrix}$

$\begin{bmatrix} 160 & 213 \end{bmatrix}$

Precision: 0.99, Recall: 0.57, F1-score: 0.72

iii. Oversampled





Best parameters for Decision Tree: {'min_samples_split': 2, 'min_samples_leaf': 1, 'max_depth': None, 'criterion': 'entropy'}

Confusion Matrix for Decision Tree:

[[2 2]

[151 222]]

Precision: 0.99, Recall: 0.60, F1-score: 0.74

Best parameters for Random Forest: {'n_estimators': 100, 'min_samples_split': 2, 'min_samples_leaf': 1, 'max_depth': None, 'bootstrap': False}

Confusion Matrix for Random Forest:

[[4 0]

[199 174]]

Precision: 1.00, Recall: 0.47, F1-score: 0.64

Best parameters for SVM: {'kernel': 'rbf', 'gamma': 'scale', 'degree': 2, 'C': 100}

Confusion Matrix for SVM:

```
[[ 0  4]
```

```
 [13 360]]
```

Precision: 0.99, Recall: 0.97, F1-score: 0.98

Best parameters for Gradient Boosting: {'subsample': 0.8, 'n_estimators': 100, 'min_samples_split': 2, 'max_depth': 10, 'learning_rate': 0.2}

Confusion Matrix for Gradient Boosting:

```
[[ 4  0]
```

```
 [318 55]]
```

Precision: 1.00, Recall: 0.15, F1-score: 0.26

Best parameters for MLP: {'solver': 'adam', 'learning_rate': 'constant', 'hidden_layer_sizes': (50,), 'alpha': 0.001, 'activation': 'relu'}

Confusion Matrix for MLP:

```
[[ 0  4]
```

```
 [12 361]]
```

Precision: 0.99, Recall: 0.97, F1-score: 0.98

Best parameters for k-NN: {'weights': 'uniform', 'p': 1, 'n_neighbors': 3}


Confusion Matrix for k-NN:

```
[[ 0  4]
```

```
 [54 319]]
```

Precision: 0.99, Recall: 0.86, F1-score: 0.92

iv. Combined



Decision Tree

Actual \ Predicted	0	1
0	2	2
1	212	161

Random Forest

Actual \ Predicted	0	1
0	2	2
1	194	179

SVM

Actual \ Predicted	0	1
0	0	4
1	32	341

Gradient Boosting

Actual \ Predicted	0	1
0	4	0
1	340	43

MLP

Actual \ Predicted	0	1
0	0	4
1	31	342

k-NN

Actual \ Predicted	0	1
0	0	4
1	78	295

Best parameters for Decision Tree: {'min_samples_split': 2, 'min_samples_leaf': 1, 'max_depth': None, 'criterion': 'entropy'}

Confusion Matrix for Decision Tree:

```
[[ 2  2]
```

```
[212 161]]
```

Precision: 0.99, Recall: 0.43, F1-score: 0.60

Best parameters for Random Forest: {'n_estimators': 100, 'min_samples_split': 2, 'min_samples_leaf': 1, 'max_depth': None, 'bootstrap': False}

Confusion Matrix for Random Forest:

```
[[ 2  2]
```

```
[194 179]]
```

Precision: 0.99, Recall: 0.48, F1-score: 0.65

Best parameters for SVM: {'kernel': 'rbf', 'gamma': 'auto', 'degree': 4, 'C': 10}

Confusion Matrix for SVM:

```
[[ 0  4]
```

```
[ 32 341]]
```

Precision: 0.99, Recall: 0.91, F1-score: 0.95

Best parameters for Gradient Boosting: {'subsample': 1.0, 'n_estimators': 200, 'min_samples_split': 10, 'max_depth': 5, 'learning_rate': 0.5}

Confusion Matrix for Gradient Boosting:

```
[[ 4  0]
```

```
[330  43]]
```

Precision: 1.00, Recall: 0.12, F1-score: 0.21

Best parameters for MLP: {'solver': 'adam', 'learning_rate': 'adaptive', 'hidden_layer_sizes': (100, 100), 'alpha': 0.01, 'activation': 'relu'}

Confusion Matrix for MLP:

```
[[ 0  4]
```

```
 [ 31 342]]
```

Precision: 0.99, Recall: 0.92, F1-score: 0.95

Best parameters for k-NN: {'weights': 'uniform', 'p': 2, 'n_neighbors': 3}

Confusion Matrix for k-NN:


```
[[ 0  4]
```

```
 [ 78 295]]
```

Precision: 0.99, Recall: 0.79, F1-score: 0.88

2. MAGIC MUSHROOMS DATASET(Balanced Dataset)

i. Original Dataset



Decision Tree

Actual \ Predicted	0	1
0	187	46
1	40	104

Random Forest

Actual \ Predicted	0	1
0	192	41
1	53	91

SVM

Actual \ Predicted	0	1
0	191	42
1	46	98

Gradient Boosting

Actual \ Predicted	0	1
0	194	39
1	52	92

MLP

Actual \ Predicted	0	1
0	194	39
1	51	93

k-NN

Actual \ Predicted	0	1
0	193	40
1	63	81

Best parameters for Decision Tree: {'min_samples_split': 2, 'min_samples_leaf': 1, 'max_depth': 3, 'criterion': 'gini'}

Confusion Matrix for Decision Tree:

[[187 46]

[40 104]]

Precision: 0.69, Recall: 0.72, F1-score: 0.71

Best parameters for Random Forest: {'n_estimators': 50, 'min_samples_split': 5, 'min_samples_leaf': 1, 'max_depth': 5, 'bootstrap': False}

Confusion Matrix for Random Forest:

[[192 41]

[53 91]]

Precision: 0.69, Recall: 0.63, F1-score: 0.66

Best parameters for SVM: {'kernel': 'linear', 'gamma': 'scale', 'degree': 2, 'C': 0.1}

Confusion Matrix for SVM:

```
[[191 42]
```

```
 [ 46 98]]
```

Precision: 0.70, Recall: 0.68, F1-score: 0.69

Best parameters for Gradient Boosting: {'subsample': 0.9, 'n_estimators': 200, 'min_samples_split': 10, 'max_depth': 3, 'learning_rate': 0.01}

Confusion Matrix for Gradient Boosting:

```
[[194 39]
```

```
 [ 52 92]]
```

Precision: 0.70, Recall: 0.64, F1-score: 0.67

Best parameters for MLP: {'solver': 'sgd', 'learning_rate': 'constant', 'hidden_layer_sizes': (100,), 'alpha': 0.001, 'activation': 'relu'}

Confusion Matrix for MLP:

```
[[194 39]
```

```
 [ 51 93]]
```

Precision: 0.70, Recall: 0.65, F1-score: 0.67

Best parameters for k-NN: {'weights': 'uniform', 'p': 1, 'n_neighbors': 10}


Confusion Matrix for k-NN:

```
[[193 40]
```

```
 [ 63 81]]
```

Precision: 0.67, Recall: 0.56, F1-score: 0.61

ii. Undersampled



Decision Tree

Actual	0	156	77
	1	29	115
		0	1

Predicted

Random Forest

Actual	0	166	67
	1	30	114
		0	1

Predicted

SVM

Actual	0	160	73
	1	31	113
		0	1

Predicted

Gradient Boosting

Actual	0	166	67
	1	34	110
		0	1

Predicted

MLP

Actual	0	164	69
	1	22	122
		0	1

Predicted

k-NN

Actual	0	163	70
	1	27	117
		0	1

Predicted

Best parameters for Decision Tree: {'min_samples_split': 2, 'min_samples_leaf': 1, 'max_depth': 3, 'criterion': 'gini'}

Confusion Matrix for Decision Tree:

[[156 77]

[29 115]]

Precision: 0.60, Recall: 0.80, F1-score: 0.68

Best parameters for Random Forest: {'n_estimators': 100, 'min_samples_split': 2, 'min_samples_leaf': 4, 'max_depth': 10, 'bootstrap': True}

Confusion Matrix for Random Forest:

[[166 67]

[30 114]]

Precision: 0.63, Recall: 0.79, F1-score: 0.70

Best parameters for SVM: {'kernel': 'rbf', 'gamma': 'scale', 'degree': 3, 'C': 1}

Confusion Matrix for SVM:

[[160 73]

[31 113]]

Precision: 0.61, Recall: 0.78, F1-score: 0.68

Best parameters for Gradient Boosting: {'subsample': 0.8, 'n_estimators': 50, 'min_samples_split': 2, 'max_depth': 5, 'learning_rate': 0.1}

Confusion Matrix for Gradient Boosting:

[[166 67]

[34 110]]

Precision: 0.62, Recall: 0.76, F1-score: 0.69

Best parameters for MLP: {'solver': 'adam', 'learning_rate': 'adaptive', 'hidden_layer_sizes': (50, 50), 'alpha': 0.0001, 'activation': 'tanh'}

Confusion Matrix for MLP:

[[164 69]

[22 122]]

Precision: 0.64, Recall: 0.85, F1-score: 0.73

Best parameters for k-NN: {'weights': 'uniform', 'p': 2, 'n_neighbors': 20}


Confusion Matrix for k-NN:

[[163 70]

[27 117]]

Precision: 0.63, Recall: 0.81, F1-score: 0.71

iii. Oversampled



Decision Tree

Actual \ Predicted	0	1
0	80	153
1	8	136

Random Forest

Actual \ Predicted	0	1
0	132	101
1	25	119

SVM

Actual \ Predicted	0	1
0	170	63
1	54	90

Gradient Boosting

Actual \ Predicted	0	1
0	112	121
1	18	126

MLP

Actual \ Predicted	0	1
0	166	67
1	43	101

k-NN

Actual \ Predicted	0	1
0	151	82
1	33	111

Best parameters for Decision Tree: {'min_samples_split': 10, 'min_samples_leaf': 1, 'max_depth': 5, 'criterion': 'entropy'}

Confusion Matrix for Decision Tree:

[[80 153]

[8 136]]

Precision: 0.47, Recall: 0.94, F1-score: 0.63

Best parameters for Random Forest: {'n_estimators': 100, 'min_samples_split': 2, 'min_samples_leaf': 1, 'max_depth': None, 'bootstrap': False}

Confusion Matrix for Random Forest:

[[132 101]

[25 119]]

Precision: 0.54, Recall: 0.83, F1-score: 0.65

Best parameters for SVM: {'kernel': 'rbf', 'gamma': 'auto', 'degree': 2, 'C': 100}

Confusion Matrix for SVM:

[[170 63]

[54 90]]

Precision: 0.59, Recall: 0.62, F1-score: 0.61

Best parameters for Gradient Boosting: {'subsample': 0.8, 'n_estimators': 100, 'min_samples_split': 2, 'max_depth': 10, 'learning_rate': 0.2}

Confusion Matrix for Gradient Boosting:

[[112 121]

[18 126]]

Precision: 0.51, Recall: 0.88, F1-score: 0.64

Best parameters for MLP: {'solver': 'adam', 'learning_rate': 'adaptive', 'hidden_layer_sizes': (50, 50), 'alpha': 0.0001, 'activation': 'tanh'}

Confusion Matrix for MLP:

[[166 67]

[43 101]]

Precision: 0.60, Recall: 0.70, F1-score: 0.65

Best parameters for k-NN: {'weights': 'distance', 'p': 1, 'n_neighbors': 5}

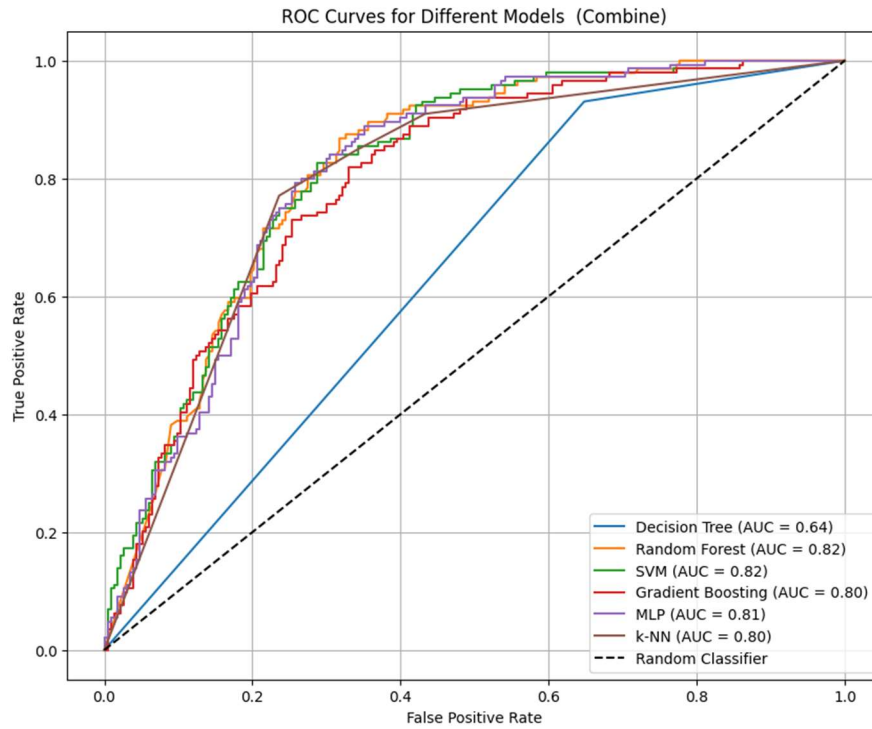
Confusion Matrix for k-NN:

[[151 82]

[33 111]]

Precision: 0.58, Recall: 0.77, F1-score: 0.66

iv. Combined



Decision Tree

Actual \ Predicted	0	1
0	82	151
1	10	134

Random Forest

Actual \ Predicted	0	1
0	135	98
1	11	133

SVM

Actual \ Predicted	0	1
0	153	80
1	21	123

Gradient Boosting

Actual \ Predicted	0	1
0	118	115
1	9	135

MLP

Actual \ Predicted	0	1
0	153	80
1	18	126

k-NN

Actual \ Predicted	0	1
0	154	79
1	22	122

Best parameters for Decision Tree: {'min_samples_split': 2, 'min_samples_leaf': 1, 'max_depth': None, 'criterion': 'entropy'}

Confusion Matrix for Decision Tree:

[[82 151]

[10 134]]

Precision: 0.47, Recall: 0.93, F1-score: 0.62

Best parameters for Random Forest: {'n_estimators': 50, 'min_samples_split': 5, 'min_samples_leaf': 2, 'max_depth': None, 'bootstrap': False}

Confusion Matrix for Random Forest:

[[135 98]

[11 133]]

Precision: 0.58, Recall: 0.92, F1-score: 0.71

Best parameters for SVM: {'kernel': 'rbf', 'gamma': 'scale', 'degree': 3, 'C': 1}

Confusion Matrix for SVM:

[[153 80]

[21 123]]

Precision: 0.61, Recall: 0.85, F1-score: 0.71

Best parameters for Gradient Boosting: {'subsample': 0.9, 'n_estimators': 200, 'min_samples_split': 2, 'max_depth': 5, 'learning_rate': 0.2}

Confusion Matrix for Gradient Boosting:

[[118 115]

[9 135]]

Precision: 0.54, Recall: 0.94, F1-score: 0.69

Best parameters for MLP: {'solver': 'adam', 'learning_rate': 'constant', 'hidden_layer_sizes': (50,),'alpha': 0.01, 'activation': 'relu'}

Confusion Matrix for MLP:

[[153 80]

[18 126]]

Precision: 0.61, Recall: 0.88, F1-score: 0.72

Best parameters for k-NN: {'weights': 'uniform', 'p': 1, 'n_neighbors': 3}

Confusion Matrix for k-NN:

[[154 79]

[22 122]]

Precision: 0.61, Recall: 0.85, F1-score: 0.71
